

Student Name: Tushar Panchal

En. No. : 21162101014

Branch: CBA

Batch: 71

Subject: ML (Machine Learning)

EXPERIMENT 2

Feature extraction using visualization and statistical test

Import dataset available at following url:

<https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data>

```
In [31]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats

# Load the dataset
path = "https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85."
df = pd.read_csv(path, na_values="?", header=None)

# Define the column headers
headers = [
    "symboling",
    "normalized-losses",
    "make",
    "fuel-type",
    "aspiration",
    "num-of-doors",
    "body-style",
    "drive-wheels",
    "engine-location",
    "wheel-base",
    "length",
    "width",
    # Added missing comma here
```

```

    "height",
    "curb-weight",
    "engine-type",
    "num-of-cylinders",
    "engine-size",
    "fuel-system",
    "bore",
    "stroke",
    "compression-ratio",
    "horsepower",
    "peak-rpm",
    "city-mpg",
    "highway-mpg",
    "price",
]

# Assign column names to the DataFrame
df.columns = headers

# Display the first 5 rows of the DataFrame
print("The first 5 rows of the dataframe:")
print(df.head())

```

The first 5 rows of the dataframe:

	symboling	normalized-losses	make	fuel-type	aspiration	\
0	3	NaN	alfa-romero	gas	std	
1	3	NaN	alfa-romero	gas	std	
2	1	NaN	alfa-romero	gas	std	
3	2	164.0	audi	gas	std	
4	2	164.0	audi	gas	std	

	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	\
0	two	convertible	rwd	front	88.6	...	
1	two	convertible	rwd	front	88.6	...	
2	two	hatchback	rwd	front	94.5	...	
3	four	sedan	fwd	front	99.8	...	
4	four	sedan	4wd	front	99.4	...	

	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	\
0	130	mpfi	3.47	2.68	9.0	111.0	
1	130	mpfi	3.47	2.68	9.0	111.0	
2	152	mpfi	2.68	3.47	9.0	154.0	
3	109	mpfi	3.19	3.40	10.0	102.0	
4	136	mpfi	3.19	3.40	8.0	115.0	

	peak-rpm	city-mpg	highway-mpg	price
0	5000.0	21	27	13495.0
1	5000.0	21	27	16500.0
2	5000.0	19	26	16500.0
3	5500.0	24	30	13950.0
4	5500.0	18	22	17450.0

[5 rows x 26 columns]

Data Preprocessing:

```

In [32]: # Replace missing values with the mean for numerical columns and the mode for categor
df['normalized-losses'].replace('?', np.nan, inplace=True)
df['bore'].replace('?', np.nan, inplace=True)
df['stroke'].replace('?', np.nan, inplace=True)
df['horsepower'].replace('?', np.nan, inplace=True)
df['peak-rpm'].replace('?', np.nan, inplace=True)
df['num-of-doors'].replace('?', np.nan, inplace=True)

```

```
# Fill missing values
df ['normalized-losses'].fillna(df['normalized-losses'].astype('float').mean(), inplace=True)
df ['bore'].fillna (df ['bore'].astype('float').mean(), inplace=True)
df ['stroke'].fillna (df ['stroke'].astype('float').mean(), inplace=True)
df ['horsepower'].fillna(df['horsepower'].astype('float').mean(), inplace=True)
df ['peak-rpm'].fillna (df ['peak-rpm'].astype('float').mean(), inplace=True)
df ['num-of-doors'].fillna (df['num-of-doors'].mode() [0], inplace=True)

#Convert data types to appropriate formats
df['price'] = df ['price'].replace('?', np.nan).astype('float')
df.dropna (subset=['price'], inplace=True) # Remove rows with NaN values in 'price'
df ['price'] = df['price'].astype('float')
df ['normalized-losses'] = df ['normalized-losses'].astype('float')
df['bore'] = df ['bore'].astype('float')
df['stroke'] = df ['stroke'].astype('float')
df ['horsepower'] = df ['horsepower'].astype('float')
df['peak-rpm'] = df['peak-rpm'].astype('float')
```

1.List down all the continuous attributes in the dataset:

```
In [33]: # Select continuous attributes (numerical columns)
continuous_attributes = df.select_dtypes(include=['float64', 'int64']).columns.tolist()

print("\nContinuous Attributes: \n")
for attribute in continuous_attributes:
    print(f'★ {attribute}')
```

Continuous Attributes:

```
★ symboling
★ normalized-losses
★ wheel-base
★ length
★ width
★ height
★ curb-weight
★ engine-size
★ bore
★ stroke
★ compression-ratio
★ horsepower
★ peak-rpm
★ city-mpg
★ highway-mpg
★ price
```

2. List down all the categorical attributes in the dataset:

```
In [34]: # Select categorical attributes (columns with object type)
categorical_attributes = df.select_dtypes(include=['object']).columns.tolist()

print("\nCategorical Attributes: \n")
for attribute in categorical_attributes:
    print(f'☀ {attribute}')
```

Categorical Attributes:

- ⚡ make
- ⚡ fuel-type
- ⚡ aspiration
- ⚡ num-of-doors
- ⚡ body-style
- ⚡ drive-wheels
- ⚡ engine-location
- ⚡ engine-type
- ⚡ num-of-cylinders
- ⚡ fuel-system

3. Draw regplot between each continuous attribute and price and write down whether that attribute is related to price or not.

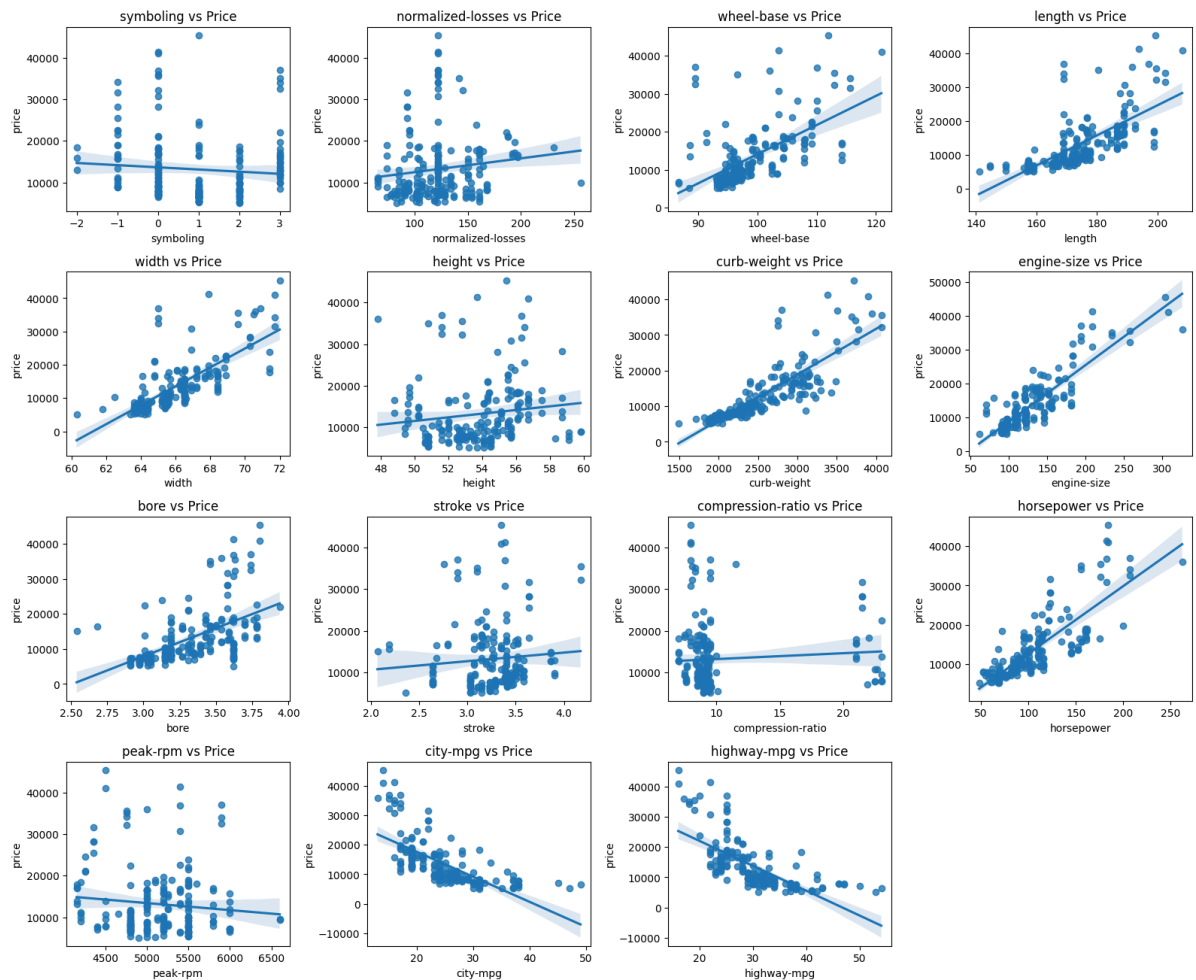
```
In [35]: # Data visualization for continuous attributes
plt.figure(figsize=(16, 16))

for i, attr in enumerate(continuous_attributes):
    if attr != 'price':
        plt.subplot(5, 4, i + 1)
        sns.regplot(x=attr, y='price', data=df)
        plt.title(f'{attr} vs Price')

plt.tight_layout()
plt.show()

# Check if continuous attributes are related to price
print("\nRelationship of Continuous Attributes with Price:\n")

for attr in continuous_attributes:
    if attr != 'price':
        correlation = df[[attr, 'price']].corr().iloc[0, 1]
        print(f'{attr}: {"✅" if abs(correlation) > 0.5 else "❌"}')
```



Relationship of Continuous Attributes with Price:

symboling: ✗
 normalized-losses: ✗
 wheel-base: ✓
 length: ✓
 width: ✓
 height: ✗
 curb-weight: ✓
 engine-size: ✓
 bore: ✓
 stroke: ✗
 compression-ratio: ✗
 horsepower: ✓
 peak-rpm: ✗
 city-mpg: ✓
 highway-mpg: ✓

4. Draw boxplot between each categorical attribute and price and write down whether that attribute is related to price or not.

```
In [36]: # Data visualization for categorical attributes
plt.figure(figsize=(22, 32))

for i, attr in enumerate(categorical_attributes):
    plt.subplot(5, 2, i + 1)
    sns.boxplot(x=attr, y='price', data=df)
    plt.title(f'{attr} vs Price')

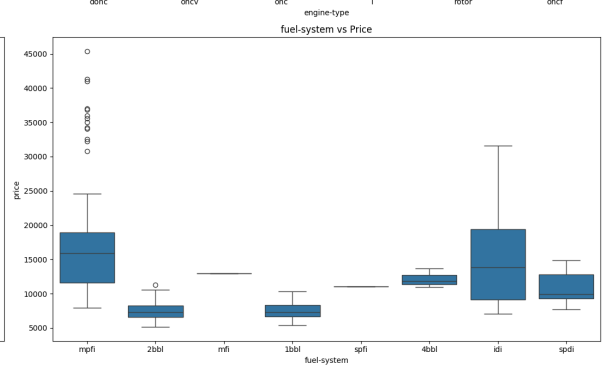


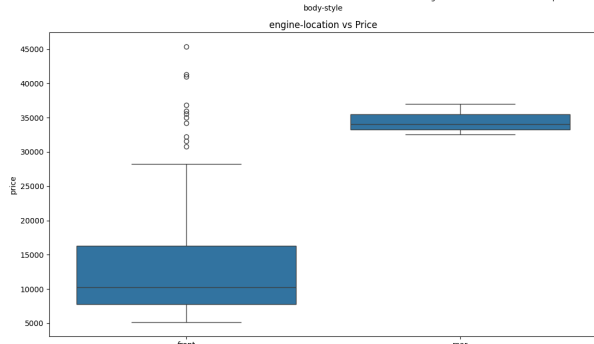

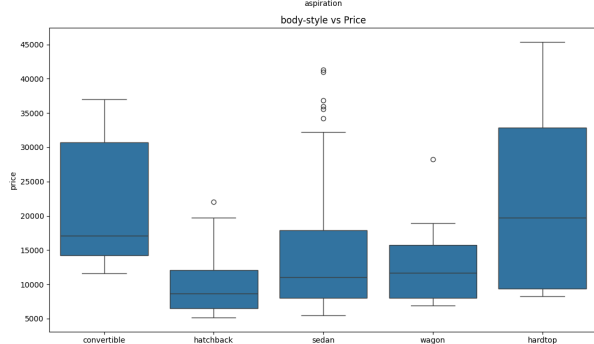


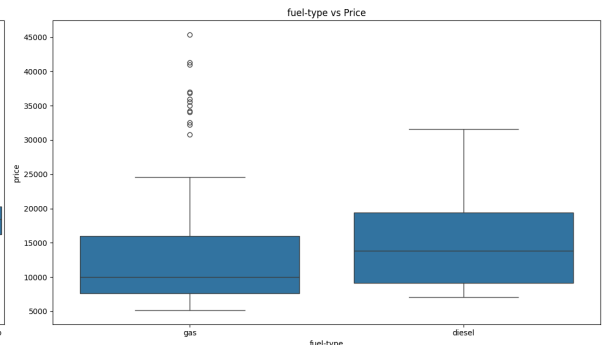
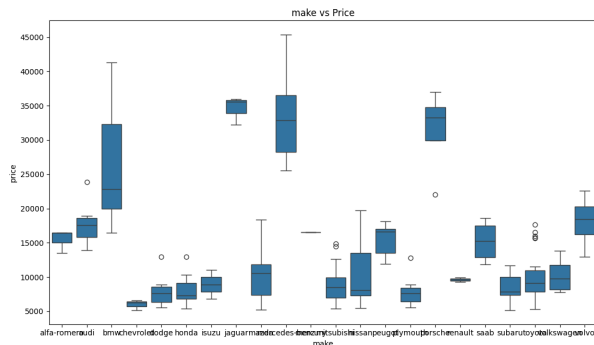
plt.tight_layout()
```

```
plt.show()











# Check if categorical attributes are related to price
print("\nRelationship of Categorical Attributes with Price:\n")

for attr in categorical_attributes:
    grouped_test = df[[attr, 'price']].groupby([attr])
    unique_values = df[attr].unique()

    if len(unique_values) > 1:
        f_val, p_val = stats.f_oneway(*[grouped_test.get_group(val)['price'] for val in unique_values])
        print(f'{attr}: {"✅" if p_val < 0.05 else "❌"}')
```



Relationship of Categorical Attributes with Price:
















make: 
fuel-type: 
aspiration: 
num-of-doors: 
body-style: 
drive-wheels: 
engine-location: 
engine-type: 
num-of-cylinders: 
fuel-system: 

5. Calculate pearson correlation between each continuous attribute and price and write down whether that attribute is related to price or not.

```
In [37]: # Print Pearson Correlation Coefficients with Price (and if related)
print("\nPearson Correlation Coefficients with Price (and if related):\n")

for attr in continuous_attributes:
    if attr != 'price':
        correlation = df[[attr, 'price']].corr().iloc[0, 1]
        print(f'{attr}: {correlation:.2f} {"(✓)" if abs(correlation) > 0.5 else " "}
```

Pearson Correlation Coefficients with Price (and if related):

symboling: -0.08 ()
normalized-losses: 0.13 ()
wheel-base: 0.58 ()
length: 0.69 ()
width: 0.75 ()
height: 0.14 ()
curb-weight: 0.83 ()
engine-size: 0.87 ()
bore: 0.54 ()
stroke: 0.08 ()
compression-ratio: 0.07 ()
horsepower: 0.81 ()
peak-rpm: -0.10 ()
city-mpg: -0.69 ()
highway-mpg: -0.70 ()

6. Calculate ANOVA between each categorical attribute and price and write down whether that attribute is related to price or not.

```
In [38]: # Print ANOVA results for categorical attributes
print("\nANOVA Results for Categorical Attributes:\n")

for attr in categorical_attributes:
    grouped_test = df[[attr, 'price']].groupby([attr])
    unique_values = df[attr].unique()

    if len(unique_values) > 1:
        f_val, p_val = stats.f_oneway(*[grouped_test.get_group(val)['price'] for val in unique_values])
        print(f'{attr}: \n\tF={f_val:.2f}, \n\tp={p_val:.2e} {"(✓)" if p_val < 0.05 else " "}
```


ANOVA Results for Categorical Attributes:

make:

F=33.23,
p=1.07e-50 (✓)

fuel-type:

F=2.45,
p=1.19e-01 (✗)

aspiration:

F=6.63,
p=1.07e-02 (✓)

num-of-doors:

F=0.36,
p=5.50e-01 (✗)

body-style:

F=9.13,
p=8.78e-07 (✓)

drive-wheels:

F=67.95,
p=3.39e-23 (✓)

engine-location:

F=24.50,
p=1.58e-06 (✓)

engine-type:

F=9.85,
p=2.09e-08 (✓)

num-of-cylinders:

F=54.94,
p=2.87e-39 (✓)

fuel-system:

F=15.02,
p=1.31e-15 (✓)

7. List down the attributes which has significant impact on price.

```
In [39]: # Lists to store attributes with significant impact on price
significant_continuous = []
significant_categorical = []

# Check for significant continuous attributes
for attr in continuous_attributes:
    if attr != 'price':
        correlation = df[[attr, 'price']].corr().iloc[0, 1]
        if abs(correlation) > 0.5:
            significant_continuous.append(attr)

# Check for significant categorical attributes
for attr in categorical_attributes:
    grouped_test = df[[attr, 'price']].groupby([attr])
    unique_values = df[attr].unique()
    mean_prices = [grouped_test.get_group(val)['price'].mean() for val in unique_values]
```

```

    if mean_prices:
        max_diff = max(mean_prices) - min(mean_prices)
        if max_diff / df['price'].mean() > 0.2: # Arbitrary threshold of 20% mean
            significant_categorical.append(attr)

# Print attributes with significant impact on price
print("\nAttributes with Significant Impact on Price:\n")

print("Continuous:\n")
for attr in significant_continuous:
    print(f'💎 {attr}')

print("\nCategorical:\n")
for attr in significant_categorical:
    print(f'🌟 {attr}')

```

Attributes with Significant Impact on Price:

Continuous:

💎 wheel-base
 💎 length
 💎 width
 💎 curb-weight
 💎 engine-size
 💎 bore
 💎 horsepower
 💎 city-mpg
 💎 highway-mpg

Categorical:

🌟 make
 🌟 fuel-type
 🌟 aspiration
 🌟 body-style
 🌟 drive-wheels
 🌟 engine-location
 🌟 engine-type
 🌟 num-of-cylinders
 🌟 fuel-system

8. Clean the assigned dataset and find important features from it.

```

In [40]: # Combine significant continuous and categorical features
important_features = significant_continuous + significant_categorical

# Print important features
print("\nImportant Features:\n")
for feature in important_features:
    print(f'🌟 {feature}')

```

Important Features:

- ★ wheel-base
- ★ length
- ★ width
- ★ curb-weight
- ★ engine-size
- ★ bore
- ★ horsepower
- ★ city-mpg
- ★ highway-mpg
- ★ make
- ★ fuel-type
- ★ aspiration
- ★ body-style
- ★ drive-wheels
- ★ engine-location
- ★ engine-type
- ★ num-of-cylinders
- ★ fuel-system