

Student Name: Tushar Panchal

En. No. : 21162101014

Branch: CBA

Batch: 71

Subject: ML (Machine Learning)

EXPERIMENT 3 and 4

Implement linear regression for given dataset and find model which has highest r^2 score and minimum MSE

Instructions:

Understand the problem statement properly Clean dataset assigned to you List of important attributes with proper justification Read sample linear regression code Answer following questions in a pdf file:

1. List down all the important attributes in the dataset
2. Write down the models you have compared.
3. Write down the model which has highest r^2 score and minimum MSE

```
In [1]: from google.colab import drive  
drive.mount('/content/drive')
```

Mounted at /content/drive

IMPORT DATASETS

```
In [2]: from os import path  
import pandas as pd  
import numpy as np  
  
#load bike data  
path_bike = "/content/drive/MyDrive/ML_DATASETS/ML_3_4/bikedata_daywise.csv"  
df_bike = pd.read_csv(path_bike, header=0)
```

```
#Load Metro Data
path_metro = "/content/drive/MyDrive/Colab Notebooks/ML-3-4/metro.csv"
df_metro = pd.read_csv(path_metro,header=0)

#Load Autos Data
path_autos = "/content/drive/MyDrive/Colab Notebooks/ML-3-4/autos.csv"
df_autos = pd.read_csv(path_autos,header=0)

#Display the first five rows of each Dataframe to verify the loading process
print("☆☆Bike Data☆☆")
print(df_bike.head(),"\n")

print("☆☆Metro Data☆☆")
print(df_metro.head(),"\n")

print("☆☆Autos Data☆☆")
print(df_autos.head(),"\n")
```

★★ Bike Data ★★

	instant	dteday	season	yr	mnth	holiday	weekday	workingday	\
0	1	2011-01-01	1	0	1	0	6	0	
1	2	2011-01-02	1	0	1	0	0	0	
2	3	2011-01-03	1	0	1	0	1	1	
3	4	2011-01-04	1	0	1	0	2	1	
4	5	2011-01-05	1	0	1	0	3	1	

	weathersit	temp	atemp	hum	windspeed	casual	registered	\
0	2	0.344167	0.363625	0.805833	0.160446	331	654	
1	2	0.363478	0.353739	0.696087	0.248539	131	670	
2	1	0.196364	0.189405	0.437273	0.248309	120	1229	
3	1	0.200000	0.212122	0.590435	0.160296	108	1454	
4	1	0.226957	0.229270	0.436957	0.186900	82	1518	

	cnt
0	985
1	801
2	1349
3	1562
4	1600

★★ Metro Data ★★

	holiday	temp	rain_1h	snow_1h	clouds_all	weather_main	\
0	NaN	288.28	0.0	0.0	40	Clouds	
1	NaN	289.36	0.0	0.0	75	Clouds	
2	NaN	289.58	0.0	0.0	90	Clouds	
3	NaN	290.13	0.0	0.0	90	Clouds	
4	NaN	291.14	0.0	0.0	75	Clouds	

	weather_description	date_time	traffic_volume
0	scattered clouds	2012-10-02 09:00:00	5545
1	broken clouds	2012-10-02 10:00:00	4516
2	overcast clouds	2012-10-02 11:00:00	4767
3	overcast clouds	2012-10-02 12:00:00	5026
4	broken clouds	2012-10-02 13:00:00	4918

★★ Autos Data ★★

	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
0	NaN	alfa-romero	gas	std	two	
1	NaN	alfa-romero	gas	std	two	
2	NaN	alfa-romero	gas	std	two	
3	164.0	audi	gas	std	four	
4	164.0	audi	gas	std	four	

	body-style	drive-wheels	engine-location	wheel-base	length	...	\
0	convertible	rwd	front	88.6	168.8	...	
1	convertible	rwd	front	88.6	168.8	...	
2	hatchback	rwd	front	94.5	171.2	...	
3	sedan	fwd	front	99.8	176.6	...	
4	sedan	4wd	front	99.4	176.6	...	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111.0	5000.0	21	
1	mpfi	3.47	2.68	9.0	111.0	5000.0	21	
2	mpfi	2.68	3.47	9.0	154.0	5000.0	19	
3	mpfi	3.19	3.40	10.0	102.0	5500.0	24	
4	mpfi	3.19	3.40	8.0	115.0	5500.0	18	

	highway-mpg	price	symboling
0	27	13495.0	3
1	27	16500.0	3
2	26	16500.0	1
3	30	13950.0	2

4 22 17450.0 2

[5 rows x 26 columns]

Clean Data

```
In [3]: def clean_data(df):
        #Convert data columns to determine if present
        if 'dteday' in df.columns:
            df['dteday'] = pd.to_datetime(df['dteday'])
        if 'data_time' in df.columns:
            df['data_time'] = pd.to_datetime(df['data_time'])

        #Replace Missing values with column means for numeric columns
        numeric_columns = df.select_dtypes(include=[np.number]).columns
        df[numeric_columns] = df[numeric_columns].fillna(df[numeric_columns].mean())

        #Replace missing values in categorical columns with the mode
        categorical_columns = df.select_dtypes(include=['object']).columns
        for col in categorical_columns:
            df[col].fillna(df[col].mode()[0], inplace=True)

        return df

#Clean each dataset
df_bike_clean = clean_data(df_bike)
df_metro_clean = clean_data(df_metro)
df_autos_clean = clean_data(df_autos)

#display the first few rows of each cleaned Dataframe to verify the cleaning process
print("★ ★ Cleaned Bike Data ★ ★")
print(df_bike_clean.head(), "\n")

print("★ ★ Cleaned Metro Data ★ ★")
print(df_metro_clean.head(), "\n")

print("★ ★ Cleaned Autos Data ★ ★")
print(df_autos_clean.head(), "\n")
```

★ ★ Cleaned Bike Data ★ ★

	instant	dteday	season	yr	mnth	holiday	weekday	workingday	\
0	1	2011-01-01	1	0	1	0	6	0	
1	2	2011-01-02	1	0	1	0	0	0	
2	3	2011-01-03	1	0	1	0	1	1	
3	4	2011-01-04	1	0	1	0	2	1	
4	5	2011-01-05	1	0	1	0	3	1	

	weathersit	temp	atemp	hum	windspeed	casual	registered	\
0	2	0.344167	0.363625	0.805833	0.160446	331	654	
1	2	0.363478	0.353739	0.696087	0.248539	131	670	
2	1	0.196364	0.189405	0.437273	0.248309	120	1229	
3	1	0.200000	0.212122	0.590435	0.160296	108	1454	
4	1	0.226957	0.229270	0.436957	0.186900	82	1518	

	cnt
0	985
1	801
2	1349
3	1562
4	1600

★ ★ Cleaned Metro Data ★ ★

	holiday	temp	rain_1h	snow_1h	clouds_all	weather_main	\
0	Labor Day	288.28	0.0	0.0	40	Clouds	
1	Labor Day	289.36	0.0	0.0	75	Clouds	
2	Labor Day	289.58	0.0	0.0	90	Clouds	
3	Labor Day	290.13	0.0	0.0	90	Clouds	
4	Labor Day	291.14	0.0	0.0	75	Clouds	

	weather_description	date_time	traffic_volume
0	scattered clouds	2012-10-02 09:00:00	5545
1	broken clouds	2012-10-02 10:00:00	4516
2	overcast clouds	2012-10-02 11:00:00	4767
3	overcast clouds	2012-10-02 12:00:00	5026
4	broken clouds	2012-10-02 13:00:00	4918

★ ★ Cleaned Autos Data ★ ★

	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
0	122.0	alfa-romero	gas	std	two	
1	122.0	alfa-romero	gas	std	two	
2	122.0	alfa-romero	gas	std	two	
3	164.0	audi	gas	std	four	
4	164.0	audi	gas	std	four	

	body-style	drive-wheels	engine-location	wheel-base	length	...	\
0	convertible	rwd	front	88.6	168.8	...	
1	convertible	rwd	front	88.6	168.8	...	
2	hatchback	rwd	front	94.5	171.2	...	
3	sedan	fwd	front	99.8	176.6	...	
4	sedan	4wd	front	99.4	176.6	...	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111.0	5000.0	21	
1	mpfi	3.47	2.68	9.0	111.0	5000.0	21	
2	mpfi	2.68	3.47	9.0	154.0	5000.0	19	
3	mpfi	3.19	3.40	10.0	102.0	5500.0	24	
4	mpfi	3.19	3.40	8.0	115.0	5500.0	18	

	highway-mpg	price	symboling
0	27	13495.0	3
1	27	16500.0	3
2	26	16500.0	1
3	30	13950.0	2

```
4          22  17450.0          2
```

```
[5 rows x 26 columns]
```

<ipython-input-3-8f957a9250ff>:15: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df[col].fillna(df[col].mode()[0],inplace=True)
```

<ipython-input-3-8f957a9250ff>:15: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df[col].fillna(df[col].mode()[0],inplace=True)
```

LinearRegression-Model for Bike Data

```
In [4]: from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

#prepare the data for regression
x_bike= df_bike_clean[['temp', 'atemp', 'hum', 'windspeed']] #Features selected
y_bike= df_bike_clean['cnt'] #Target Variable

#Split the data into training and testing sets
x_train_bike,x_test_bike,y_train_bike,y_test_bike = train_test_split(x_bike,y_bike,

lr_bike = LinearRegression()

#Train the Model on the training data
lr_bike.fit(x_train_bike,y_train_bike)

#Make Prediction on the test data
y_pred_bike=lr_bike.predict(x_test_bike)

#Evaluate the model using R^2 score
r2_bike=r2_score(y_test_bike,y_pred_bike)
print(f"Linear Regression R^2 Score for Bike Data: {r2_bike:.4f}")

#Perform cross-validation and calculate the mean cross-validation score
cv_scores_bike = cross_val_score(lr_bike,x_bike,y_bike,cv=5)
cv_mean_bike = np.mean(cv_scores_bike)
print(f"Linear Regression Cross-Validation Score for Bike Data: {cv_mean_bike:.4f}")
```

```
Linear Regression R^2 Score for Bike Data: 0.4995
```

```
Linear Regression Cross-Validation Score for Bike Data: -1.9649
```

Ridge-Model on Bike Data

```
In [5]: from sklearn.linear_model import Ridge

ridge_bike=Ridge(alpha=1.0)

#Train the model on the training data
ridge_bike.fit(x_train_bike,y_train_bike)

#Make prediction on the test data
y_pred_ridge_bike=ridge_bike.predict(x_test_bike)

#Evaluate the model using R^2 score
r2_ridge_bike=r2_score(y_test_bike,y_pred_ridge_bike)
print(f"Ridge Regression R^2 Score for Bike Data: {r2_ridge_bike:.4f}")

#Perform cross-validation and calculate mean cross-validation score
cv_scores_ridge_bike=cross_val_score(ridge_bike,x_bike,y_bike,cv=5)
cv_mean_ridge_bike=np.mean(cv_scores_ridge_bike)
print(f"Ridge Regression Cross-Validation Score for Bike Data: {cv_mean_ridge_bike:.4f}")
```

Ridge Regression R² Score for Bike Data: 0.4869

Ridge Regression Cross-Validation Score for Bike Data: -1.8712

Lasso-Model on Bike Data

```
In [7]: from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Lasso
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import r2_score
import numpy as np

# Standardize the features
scaler = StandardScaler()
X_bike_scaled = scaler.fit_transform(x_bike)

# Split the scaled data into training and testing sets
X_train_bike_scaled, X_test_bike_scaled, y_train_bike, y_test_bike = train_test_split(X_bike_scaled, y_bike, test_size=0.2, random_state=42)

# Initialize Lasso with increased max_iter
lasso_bike = Lasso(alpha=0.1, max_iter=10000)

# Train the model on the scaled training data
lasso_bike.fit(X_train_bike_scaled, y_train_bike)

# Make predictions on the scaled test data
y_pred_lasso_bike = lasso_bike.predict(X_test_bike_scaled)

# Evaluate the model using R^2 score
r2_lasso_bike = r2_score(y_test_bike, y_pred_lasso_bike)
print(f"Lasso Regression R^2 Score for Bike Dataset (Scaled): {r2_lasso_bike:.4f}")

# Perform cross-validation and calculate the mean cross-validation score
cv_scores_lasso_bike = cross_val_score(lasso_bike, X_bike_scaled, y_bike, cv=5)
cv_mean_lasso_bike = np.mean(cv_scores_lasso_bike)
print(f"Lasso Regression Cross-Validation Score for Bike Dataset (Scaled): {cv_mean_lasso_bike:.4f}")
```

Lasso Regression R² Score for Bike Dataset (Scaled): 0.4994

Lasso Regression Cross-Validation Score for Bike Dataset (Scaled): -1.9648

LinearRegression-Model for Metro Data

```
In [8]: # Prepare the data for regression
X_metro = df_metro_clean[['temp', 'rain_1h', 'snow_1h', 'clouds_all']] # Features
y_metro = df_metro_clean['traffic_volume'] # Target variable

# Split the data into training and testing sets
X_train_metro, X_test_metro, y_train_metro, y_test_metro = train_test_split(X_metro, y_metro,
                                     test_size=0.2, random_state=42)

# Initialize the Linear Regression model
lr_metro = LinearRegression()

# Train the model on the training data
lr_metro.fit(X_train_metro, y_train_metro)

# Make predictions on the test data
y_pred_metro = lr_metro.predict(X_test_metro)

# Evaluate the model using R² score
r2_metro = r2_score(y_test_metro, y_pred_metro)
print(f"Linear Regression R² Score for Metro Dataset: {r2_metro:.4f}")

# Perform cross-validation and calculate the mean cross-validation score
cv_scores_metro = cross_val_score(lr_metro, X_metro, y_metro, cv=5)
cv_mean_metro = np.mean(cv_scores_metro)
print(f"Linear Regression Cross-Validation Score for Metro Dataset: {cv_mean_metro:.4f}")
```

Linear Regression R² Score for Metro Dataset: 0.0234

Linear Regression Cross-Validation Score for Metro Dataset: -2.5345

Ridge-Model on Metro Data

```
In [9]: # Initialize the Ridge Regression model
ridge_metro = Ridge(alpha=1.0)

# Train the model on the training data
ridge_metro.fit(X_train_metro, y_train_metro)

# Make predictions on the test data
y_pred_ridge_metro = ridge_metro.predict(X_test_metro)

# Evaluate the model using R² score
r2_ridge_metro = r2_score(y_test_metro, y_pred_ridge_metro)
print(f"Ridge Regression R² Score for Metro Dataset: {r2_ridge_metro:.4f}")

# Perform cross-validation and calculate the mean cross-validation score
cv_scores_ridge_metro = cross_val_score(ridge_metro, X_metro, y_metro, cv=5)
cv_mean_ridge_metro = np.mean(cv_scores_ridge_metro)
print(f"Ridge Regression Cross-Validation Score for Metro Dataset: {cv_mean_ridge_metro:.4f}")
```

Ridge Regression R² Score for Metro Dataset: 0.0234

Ridge Regression Cross-Validation Score for Metro Dataset: -2.5344

Lasso-Model on Metro Data

```
In [10]: # Feature Scaling
scaler = StandardScaler()
```



```

X_metro_scaled = scaler.fit_transform(X_metro)

# Initialize Lasso Regression model
lasso_metro = Lasso(alpha=1.0, max_iter=10000)

# Train the model on the training data
lasso_metro.fit(X_train_metro, y_train_metro)

# Make predictions on the test data
y_pred_lasso_metro = lasso_metro.predict(X_test_metro)

# Evaluate the model using R² score
r2_lasso_metro = r2_score(y_test_metro, y_pred_lasso_metro)
print(f"Lasso Regression R² Score for Metro Dataset: {r2_lasso_metro:.4f}")

# Perform cross-validation and calculate the mean cross-validation score
cv_scores_lasso_metro = cross_val_score(lasso_metro, X_metro_scaled, y_metro, cv=5)
cv_mean_lasso_metro = np.mean(cv_scores_lasso_metro)
print(f"Lasso Regression Cross-Validation Score for Metro Dataset: {cv_mean_lasso_metro:.4f}")

```

Lasso Regression R² Score for Metro Dataset: 0.0234

Lasso Regression Cross-Validation Score for Metro Dataset: -0.3194

LinearRegression-Model for Auto Data

```

In [11]: # Drop rows with missing target values
df_autos_clean = df_autos_clean.dropna(subset=['price'])

# Convert categorical variables to dummy variables
df_autos_clean = pd.get_dummies(df_autos_clean, columns=['make', 'fuel-type', 'aspiration', 'body-style', 'drive-wheels', 'engine-location', 'engine-type', 'num-of-cylinders', 'transmission'])

# Prepare the data for regression
X_autos = df_autos_clean.drop(columns=['price']) # Features
y_autos = df_autos_clean['price'] # Target variable

# Split the data into training and testing sets
X_train_autos, X_test_autos, y_train_autos, y_test_autos = train_test_split(X_autos, y_autos, test_size=0.2, random_state=42)

# Initialize Linear Regression model
lr_autos = LinearRegression()

# Train the model on the training data
lr_autos.fit(X_train_autos, y_train_autos)

# Make predictions on the test data
y_pred_autos = lr_autos.predict(X_test_autos)

# Evaluate the model using R² score
r2_autos = r2_score(y_test_autos, y_pred_autos)
print(f"Linear Regression R² Score for Autos Dataset: {r2_autos:.4f}")

# Perform cross-validation and calculate the mean cross-validation score
cv_scores_autos = cross_val_score(lr_autos, X_autos, y_autos, cv=5)
cv_mean_autos = np.mean(cv_scores_autos)
print(f"Linear Regression Cross-Validation Score for Autos Dataset: {cv_mean_autos:.4f}")

```

Linear Regression R² Score for Autos Dataset: 0.8902

Linear Regression Cross-Validation Score for Autos Dataset: -0.1857

Ridge-Model on Auto Data

```
In [12]: # Initialize Ridge Regression model
ridge_autos = Ridge(alpha=1.0)

# Train the model on the training data
ridge_autos.fit(X_train_autos, y_train_autos)

# Make predictions on the test data
y_pred_ridge_autos = ridge_autos.predict(X_test_autos)

# Evaluate the model using R² score
r2_ridge_autos = r2_score(y_test_autos, y_pred_ridge_autos)
print(f"Ridge Regression R² Score for Autos Dataset: {r2_ridge_autos:.4f}")

# Perform cross-validation and calculate the mean cross-validation score
cv_scores_ridge_autos = cross_val_score(ridge_autos, X_autos, y_autos, cv=5)
cv_mean_ridge_autos = np.mean(cv_scores_ridge_autos)
print(f"Ridge Regression Cross-Validation Score for Autos Dataset: {cv_mean_ridge_a

Ridge Regression R² Score for Autos Dataset: 0.8806
Ridge Regression Cross-Validation Score for Autos Dataset: 0.3715
```

Lasso-Model on Auto Data

```
In [13]: # Feature Scaling
scaler = StandardScaler()
X_autos_scaled = scaler.fit_transform(X_autos)

# Initialize Lasso Regression model
lasso_autos = Lasso(alpha=1.0, max_iter=10000)

# Train the model on the training data
lasso_autos.fit(X_train_autos, y_train_autos)

# Make predictions on the test data
y_pred_lasso_autos = lasso_autos.predict(X_test_autos)

# Evaluate the model using R² score
r2_lasso_autos = r2_score(y_test_autos, y_pred_lasso_autos)
print(f"Lasso Regression R² Score for Autos Dataset: {r2_lasso_autos:.4f}")

# Perform cross-validation and calculate the mean cross-validation score
cv_scores_lasso_autos = cross_val_score(lasso_autos, X_autos_scaled, y_autos, cv=5)
cv_mean_lasso_autos = np.mean(cv_scores_lasso_autos)
print(f"Lasso Regression Cross-Validation Score for Autos Dataset: {cv_mean_lasso_a

Lasso Regression R² Score for Autos Dataset: 0.8894
Lasso Regression Cross-Validation Score for Autos Dataset: 0.1326
```

Comparison of each model for Bike Data

```
In [14]: # Create a dictionary to hold the metrics
metrics_bike = {
    'Model': ['Linear Regression', 'Ridge Regression', 'Lasso Regression'],
    'R² Score': [r2_bike, r2_ridge_bike, r2_lasso_bike],
    'Cross-Validation Score': [cv_mean_bike, cv_mean_ridge_bike, cv_mean_lasso_bike]
}

# Create a DataFrame from the metrics dictionary
df_metrics_bike = pd.DataFrame(metrics_bike)
```

```
# Print the comparison
print("Bike Data Model Comparison:\n")
print(df_metrics_bike)
```

Bike Data Model Comparison:

	Model	R ² Score	Cross-Validation Score
0	Linear Regression	0.499472	-1.964897
1	Ridge Regression	0.486895	-1.871240
2	Lasso Regression	0.499442	-1.964756

Comparison of each model for Metro Data

```
In [15]: # Create a dictionary to hold the metrics
metrics_metro = {
    'Model': ['Linear Regression', 'Ridge Regression', 'Lasso Regression'],
    'R2 Score': [r2_metro, r2_ridge_metro, r2_lasso_metro],
    'Cross-Validation Score': [cv_mean_metro, cv_mean_ridge_metro, cv_mean_lasso_metro]
}

# Create a DataFrame from the metrics dictionary
df_metrics_metro = pd.DataFrame(metrics_metro)

# Print the comparison
print("Metro Data Model Comparison:")
print(df_metrics_metro)
```

Metro Data Model Comparison:

	Model	R ² Score	Cross-Validation Score
0	Linear Regression	0.023424	-2.534509
1	Ridge Regression	0.023425	-2.534377
2	Lasso Regression	0.023427	-0.319376

Comparison of each model for Auto Data

```
In [16]: # Create a dictionary to hold the metrics
metrics_autos = {
    'Model': ['Linear Regression', 'Ridge Regression', 'Lasso Regression'],
    'R2 Score': [r2_autos, r2_ridge_autos, r2_lasso_autos],
    'Cross-Validation Score': [cv_mean_autos, cv_mean_ridge_autos, cv_mean_lasso_autos]
}

# Create a DataFrame from the metrics dictionary
df_metrics_autos = pd.DataFrame(metrics_autos)

# Print the comparison
print("Autos Data Model Comparison:")
print(df_metrics_autos)
```

Autos Data Model Comparison:

	Model	R ² Score	Cross-Validation Score
0	Linear Regression	0.890177	-0.185669
1	Ridge Regression	0.880586	0.371468
2	Lasso Regression	0.889406	0.132572